

Qwen vs GLM 5.3 Controlled Benchmark Comparison

This is a descriptive comparison using the same frozen benchmark inputs. It does not combine results into a broad accuracy claim.

Metric	Qwen	GLM 5.3
QA reviewable PASS	6/6	5/6
QA INCONCLUSIVE	0/6	1/6
QA correctness (reviewed)	7.00	10.00
QA groundedness (reviewed)	8.33	10.00
Documentation reviewable	6/6	0/6
Mean latency	46.271s	8.078s
Median latency	45.810s	9.818s

Observations

- GLM produced complete reviewable outputs for five QA cases; one QA output was truncated with finish_reason=length.
- GLM Documentation encountered four invalid_response_empty_content failures and two truncated structured outputs.
- Qwen produced reviewable outputs for all twelve generation cases after replay; its quality scores are from the separate Qwen quality artifact.
- Citation validity was inherited from trusted frozen metadata and was not re-created from model text.
- Token usage was unavailable for both providers.
- This comparison does not establish universal model superiority or statistical significance.

Artifact References

- Qwen evaluation: qwen_quality_evaluation.json
- GLM results: glm_results.json
- GLM evaluation: glm_quality_evaluation.json
- Frozen benchmark cases SHA: 5fbec49e1ad4c70af6a8aafb028f473afbdeb807d189070914b778ed3e9699af
- Frozen evidence SHA: 2359b07c36f19d47faf0171de0ab5e48ebc8b2f4620b6a8a8a6865cf75cc4c83