

Controlled Bilingual Model Benchmark v1 - Qwen Baseline

1. Experiment

This report records the local-Qwen baseline. Retrieval evidence and prompts were frozen before generation. GLM/AvalAI was not called. The replay used the exact frozen LLM inputs after the initial local-provider execution failure.

Benchmark ID: codecompass_controlled_bilingual_model_comparison_v1

Benchmark cases SHA-256: 5fbec49e1ad4c70af6a8aabf028f473afbdeb807d189070914b778ed3e9699af

Frozen retrieval evidence SHA-256: 2359b07c36f19d47faf0171de0ab5e48ebc8b2f4620b6a8a8a6865cf75cc4c83

Provider/model: ollama / qwen2.5-coder-3b-codecompass:latest

Generation: temperature 0, max_tokens 1200, provider retries 0, manual reruns 0

2. Dataset and Indexes

Repository	Files	Symbols	Chunks	Vectors	IDs	Generation	Complete
hospital_system	33	409	409	409	True	True	True
cs_bookstore	56	144	144	144	True	True	True
codecompass	96	1220	1220	1220	True	True	True

All repositories were indexed once in isolated storage. No re-index occurred during replay.

3. Frozen Retrieval Evidence

Search executions: 54. QA evidence captures: 6. Documentation evidence captures: 6.

4. Qwen Generation Replay

The original 12 provider failures remain embedded in qa_results and documentation_results. The replay appended successful outputs in replay_results using the same frozen prompts, contexts, and settings. No retries were used.

Case	Type	Language	Status	Latency (s)	Output length	Finish reason	Token usage
CB-QA-H-EN	QA	en	REVIEWABLE	22.691	274	stop	اندازه‌گیری نشده
CB-QA-H-FA	QA	fa	REVIEWABLE	11.876	156	stop	اندازه‌گیری نشده
CB-QA-B-EN	QA	en	REVIEWABLE	9.167	322	stop	اندازه‌گیری نشده
CB-QA-B-FA	QA	fa	REVIEWABLE	27.032	500	stop	اندازه‌گیری نشده
CB-QA-C-EN	QA	en	REVIEWABLE	23.423	753	stop	اندازه‌گیری نشده
CB-QA-C-FA	QA	fa	REVIEWABLE	65.131	1090	stop	اندازه‌گیری نشده
CB-DOC-H-EN	Documentation	en	REVIEWABLE	45.209	754	stop	اندازه‌گیری نشده
CB-DOC-H-FA	Documentation	fa	REVIEWABLE	67.449	758	stop	اندازه‌گیری نشده
CB-DOC-B-EN	Documentation	en	REVIEWABLE	46.410	956	stop	اندازه‌گیری نشده
CB-DOC-B-FA	Documentation	fa	REVIEWABLE	48.971	555	stop	اندازه‌گیری نشده
CB-DOC-C-EN	Documentation	en	REVIEWABLE	93.381	1694	stop	اندازه‌گیری نشده
CB-DOC-C-FA	Documentation	fa	REVIEWABLE	94.510	1023	اندازه‌گیری نشده	اندازه‌گیری نشده

Replay Summary

- Successful replay outputs: 12/12
- Mean latency: 46.271s
- Median latency: 45.810s
- Range: 9.167s - 94.510s
- Token usage: اندازه‌گیری نشده؛ provider response object does not expose it.
- Semantic scores: اندازه‌گیری نشده؛ this replay did not perform human scoring.

5. Exact Captured Outputs

CB-QA-H-EN

In the Patient.__init__ method, the following attributes are assigned: - first_name - last_name - national_id - phone_num - password - gender - city - insurance_num The attribute full_name is calculated as a combination of first_name and last_name.

CB-QA-H-FA

dequeue باعث خطای "Queue is empty" است. در صف غیرخالی، ابتدا فرانت صف به یک واحد بروزرسانی شده و سپس عنصر در فرانت جدید برگردانده می‌شود. در صف خالی، متد

CB-QA-B-EN

The frozen retrieval/index evidence remains valid and unchanged. The initial local-provider failures are preserved as historical attempts; the replay produced reviewable Qwen outputs. No GLM/AvalAI request was made, so no model comparison is present in this artifact. Human quality scores, hallucination labels, and independent semantic judgments remain اندازه‌گیری نشده and must not be inferred from generated text alone.

7. Artifact Integrity

- benchmark_cases.json is frozen; its SHA-256 is recorded above.
- frozen_retrieval_evidence.json is frozen; its SHA-256 is recorded above.
- qwen_results.json contains both initial failures and replay outputs.
- No source code, index, retrieval evidence, prompt, or configuration was changed.