

Qwen Baseline Quality Evaluation

Executive Summary

This report evaluates the completed Qwen replay against frozen benchmark expectations and frozen source/citation evidence. No indexing, retrieval, prompt, or model execution occurred in this scoring phase.

Benchmark: codecompass_controlled_bilingual_model_comparison_v1

Cases hash: 5fbec49e1ad4c70af6a8aabf028f473afbdeb807d189070914b778ed3e9699af

Evidence hash: 2359b07c36f19d47faf0171de0ab5e48ebc8b2f4620b6a8a8a6865cf75cc4c83

QA Results

Case	Correctness	Groundedness	Citation	Completeness	Hallucination	Issues
CB-QA-H-EN	8/10	10/10	10/10	partially complete	none	1
CB-QA-H-FA	6/10	7/10	10/10	partially complete	major	1
CB-QA-B-EN	6/10	8/10	10/10	partially complete	minor	1
CB-QA-B-FA	9/10	9/10	10/10	complete	none	0
CB-QA-C-EN	9/10	10/10	10/10	partially complete	none	1
CB-QA-C-FA	4/10	6/10	10/10	partially complete	major	1

Average correctness: 7.00/10

Average groundedness: 8.33/10

Average citation accuracy: 10.00/10

Documentation Results

Case	Overall	Purpose	Parameters	Return	Behavior	Dependencies	Citation	Completeness	Hallucination
CB-DOC-H-EN	7/10	9	10	9	9	8	10	complete	minor
CB-DOC-H-FA	9/10	9	10	9	10	8	10	complete	none
CB-DOC-B-EN	9/10	10	10	8	10	9	10	complete	none
CB-DOC-B-FA	7/10	8	8	7	8	9	10	partially complete	none
CB-DOC-C-EN	7/10	9	9	8	8	8	10	partially complete	none
CB-DOC-C-FA	0/10	0	0	0	0	0	10	incomplete	none

Average documentation quality: 6.50/10

Known Weaknesses

- CB-QA-H-EN (medium): appointments missing Expected: all ten attributes Actual: nine listed
- CB-QA-H-FA (high): empty queue called error Expected: literal Queue is empty. Actual: error claimed
- CB-QA-B-EN (medium): super save omitted; uniqueness invented Expected: super save; no uniqueness claim Actual: both issues
- CB-QA-C-EN (low): unique sorted omitted Expected: unique sorted results Actual: generic iteration
- CB-QA-C-FA (high): garbled/wrong provider method Expected: exact service sequence Actual: _generate named; key facts omitted
- CB-DOC-H-EN (medium): unsupported side effect Expected: source-supported behavior only Actual: state update claimed
- CB-DOC-B-FA (medium): details omitted Expected: both exact facts Actual: generic save/success
- CB-DOC-C-EN (medium): details omitted Expected: all atomicity facts Actual: high-level summary
- CB-DOC-C-FA (critical): truncated invalid JSON Expected: complete object Actual: ends mid-string

Latency

Mean: 46.271s; median: 45.810s; range: 9.167s–94.510s.

Token usage: نکرد expose آن را Ollama abstraction اندازه‌گیری نشده؛

Limitations

Scores are evidence-based human-style judgments against frozen source facts. They are not automatically derived from generation success. No GLM comparison was run. The complete generated outputs remain in qwen_quality_evaluation.json.