
CODECOMPASS LOCAL

M24 Acceptance Baseline Evaluation

Retained acceptance observations before later model comparison

ACCEPTANCE_ISSUES_FOUND

1. Executive Summary

This archive preserves the retained M24 acceptance results before any later model validation. No indexing, retrieval, model, documentation, or provider request was rerun to create it. Fields absent from retained output are marked **اندازه گیری نشده** rather than reconstructed.

Final acceptance verdict: **ACCEPTANCE_ISSUES_FOUND**.

All three fresh indexes were structurally and vector complete. Search citations resolved correctly, but two semantic searches missed the expected target in the top five. The local Qwen baseline was mixed: one factual case was fully correct, several answers were incomplete or unsupported, two insufficient-evidence checks hallucinated nonexistent behavior, and two answers were truncated. Four retained OpenAI-compatible QA attempts produced no reviewable answer. The frontend was broadly functional, with one project-switch timing race.

2. Dataset Identity and Configuration

Field	Value
Dataset	m24_acceptance_baseline
Recorded	2026-09-02
CodeCompass commit	9634bdab32af0b262a7e1cae97aea0f251804f6f
Embedding	ollama / nomic-embed-text-local:latest / 768 dimensions
Local QA	ollama / qwen2.5-coder-3b-codecompass:latest
Search	lexical, semantic, hybrid

3. Indexing Results

Repository	Commit	Files	Symbols	Chunks	Vectors	Complete	Elapsed
Hospital-System	da0b1f8cf04f36d6281a5fd44b797ad195147633	33	409	409	409	True	100.465 s
CS-Bookstore	56ee148f3b4b8bc1f4eaaaf921e22f280fa14ad7a	56	144	144	144	True	55.342 s
CodeCompass	9634bdab32af0b262a7e1cae97aea0f251804f6f	96	1220	1220	1220	True	43057.811 s

CodeCompass elapsed wall time includes session interruptions and is not a comparable pure indexing latency. Every repository had exact SQLite/physical-Chroma ID equality and matching SQLite/active vector generations.

4. Search Results

Manual relevance scoring used only retained expected-target rank: rank 1 = 10, rank 2 = 8, and target absent from top 5 = 0. This is a descriptive baseline rule, not a statistical benchmark metric.

Case	Repository	Lang	Method	Target rank	Top symbol	Relevance	Citation	Failure
SEARCH-H-EN-LEX	Hospital-System	English	lexical	1	Patient.__init__	10	valid	اندازه گیری نشده
SEARCH-H-EN-SEM	Hospital-System	English	semantic	اندازه گیری نشده	HashTable.__init__	0	valid	missed target
SEARCH-H-EN-HYB	Hospital-System	English	hybrid	2	HashTable.__init__	8	valid	wrong ranking
SEARCH-H-FA-LEX	Hospital-System	Persian	lexical	1	QueueArray.dequeue	10	valid	اندازه گیری نشده

Case	Repository	Lang	Method	Target rank	Top symbol	Relevance	Citation	Failure
SEARCH-H-FA-SEM	Hospital-System	Persian	semantic	2	Queue.dequeue	8	valid	wrong ranking
SEARCH-H-FA-HYB	Hospital-System	Persian	hybrid	1	QueueArray.dequeue	10	valid	اندازه گیری نشده
SEARCH-B-EN-LEX	CS-Bookstore	English	lexical	1	Book.save	10	valid	اندازه گیری نشده
SEARCH-B-EN-SEM	CS-Bookstore	English	semantic	1	Book.save	10	valid	اندازه گیری نشده
SEARCH-B-EN-HYB	CS-Bookstore	English	hybrid	1	Book.save	10	valid	اندازه گیری نشده
SEARCH-B-FA-LEX	CS-Bookstore	Persian	lexical	1	CommentForm.clean_name	10	valid	اندازه گیری نشده
SEARCH-B-FA-SEM	CS-Bookstore	Persian	semantic	1	CommentForm.clean_name	10	valid	اندازه گیری نشده
SEARCH-B-FA-HYB	CS-Bookstore	Persian	hybrid	1	CommentForm.clean_name	10	valid	اندازه گیری نشده
SEARCH-C-EN-LEX	CodeCompass	English	lexical	1	APIRuntime.vector_generation_matches	10	valid	اندازه گیری نشده
SEARCH-C-EN-SEM	CodeCompass	English	semantic	1	APIRuntime.vector_generation_matches	10	valid	اندازه گیری نشده
SEARCH-C-EN-HYB	CodeCompass	English	hybrid	1	APIRuntime.vector_generation_matches	10	valid	اندازه گیری نشده
SEARCH-C-FA-LEX	CodeCompass	Persian	lexical	1	RepositoryIndexCoordinator.index_repository	10	valid	اندازه گیری نشده
SEARCH-C-FA-SEM	CodeCompass	Persian	semantic	اندازه گیری نشده	test_final_source_check_discards_candidate_and_preserves_previous_index_change_before_final_scan	0	valid	missed target
SEARCH-C-FA-HYB	CodeCompass	Persian	hybrid	1	RepositoryIndexCoordinator.index_repository	10	valid	اندازه گیری نشده

The complete per-result payload (all ranks, paths, symbol types, line ranges, chunk IDs, and citation records) was not retained for all executions and is therefore **اندازه گیری نشده** in the authoritative JSON. Retained target ranks, top symbols, and successful source-navigation validation are preserved.

5. Grounded QA Results

Case	Repository	Lang	Provider	Correct	Grounded	Citation	Hallucination	Completeness	Finish	Seconds
QA-H-EN	Hospital-System	English	ollama	8	10	10	none	partially complete	stop	23.014
QA-H-FA	Hospital-System	Persian	ollama	6	8	10	minor	complete	stop	11.753
QA-B-EN	CS-Bookstore	English	ollama	7	7	10	minor	complete	stop	9.325
QA-B-FA	CS-Bookstore	Persian	ollama	5	6	10	minor	partially complete	stop	27.044
QA-C-EN	CodeCompass	English	ollama	10	10	10	none	complete	stop	18.495
QA-C-FA	CodeCompass	Persian	ollama	2	4	10	major	incomplete	length	58.827
QA-H-INS	Hospital-System	English	ollama	10	10	اندازه گیری نشده	none	complete	stop	12.81
QA-B-INS	CS-Bookstore	Persian	ollama	0	0	اندازه گیری نشده	critical	incomplete	stop	33.863
QA-C-INS	CodeCompass	English	ollama	0	0	اندازه گیری نشده	critical	incomplete	length	48.915
QA-CLOUD-H-EN	Hospital-System	English	openai_compatible	اندازه گیری نشده	اندازه گیری نشده	اندازه گیری نشده	اندازه گیری نشده	اندازه گیری نشده	اندازه گیری نشده	5.309
QA-CLOUD-H-FA	Hospital-System	Persian	openai_compatible	اندازه گیری نشده	اندازه گیری نشده	اندازه گیری نشده	اندازه گیری نشده	اندازه گیری نشده	اندازه گیری نشده	14.863

Case	Repository	Lang	Provider	Correct	Grounded	Citation	Hallucination	Completeness	Finish	Seconds
QA-CLOUD-B-EN	CS-Bookstore	English	openai_compatible	اندازه گیری نشده	اندازه گیری نشده	اندازه گیری نشده	اندازه گیری نشده	اندازه گیری نشده	اندازه گیری نشده	10.915
QA-CLOUD-B-FA	CS-Bookstore	Persian	openai_compatible	اندازه گیری نشده	اندازه گیری نشده	اندازه گیری نشده	اندازه گیری نشده	اندازه گیری نشده	اندازه گیری نشده	15.518

QA-H-EN - Hospital-System

Question: What attributes are assigned in Patient.__init__?

Provider/model: ollama / qwen2.5-coder-3b-codecompass:latest

Retrieval: hybrid

Finish reason: stop

Response time: 23.014 s

Token information: اندازه گیری نشده

Exact retained answer:

In the Patient.__init__ method, the following attributes are assigned:

- first_name
- last_name
- national_id
- phone_num
- password
- gender
- city
- insurance_num

The attribute full_name is calculated as a combination of first_name and last_name.

Issue: self.appointments was omitted.

Expected: Enumerate all ten assigned attributes, including self.appointments.

Actual: Nine expected attributes were covered; self.appointments was missing.

Severity/category: medium / missing information

QA-H-FA - Hospital-System

Question: در صف خالی و غیرخالی چه رفتاری دارد؟ QueueArray.dequeue متد

Provider/model: ollama / qwen2.5-coder-3b-codecompass:latest

Retrieval: hybrid

Finish reason: stop

Response time: 11.753 s

Token information: اندازه گیری نشده

Exact retained answer:

در صف خالی، متد dequeue باعث خطای "Queue is empty" است. در صف غیرخالی، ابتدا فرانت صف به یک واحد بروزرسانی شده و سپس عنصر در فرانت جدید برگردانده می‌شود.

Issue: The empty-queue return value was described as an error.

Expected: Return the string Queue is empty. for an empty queue; otherwise increment front and return queue[front - 1].

Actual: Called the returned string an error, while correctly summarizing the non-empty branch.

Severity/category: medium / wrong interpretation

QA-B-EN - CS-Bookstore

Question: How does Book.save populate the slug before saving?

Provider/model: o11ama / qwen2.5-coder-3b-codecompass:latest

Retrieval: hybrid

Finish reason: stop

Response time: 9.325 s

Token information: اندازه گیری نشده

Exact retained answer:

Book.save populates the slug before saving by checking if the slug attribute is empty. If it is, it uses the slugify function from the slugify module to generate a slug based on the title attribute and assigns it to the slug attribute. This ensures that every book has a unique slug before being saved to the database.

Issue: The answer asserted an unsupported uniqueness guarantee and imprecise slugify provenance.

Expected: If slug is falsy, assign slugify(self.title), then call super().save(*args, **kwargs).

Actual: Core assignment was correct, but uniqueness was claimed without source evidence and the super save call was omitted.

Severity/category: medium / unsupported claim

QA-B-FA - CS-Bookstore

Question: چه نام‌هایی را رد می‌کند و در حالت معتبر چه برمی‌گرداند؟ متد CommentForm.clean_name

Provider/model: o11ama / qwen2.5-coder-3b-codecompass:latest

Retrieval: hybrid

Finish reason: stop

Response time: 27.044 s

Token information: اندازه گیری نشده

Exact retained answer:

متد CommentForm.clean_name رد می‌کند:

1. "نام معتبر نیست!" را در حالت معتبر برمی‌گرداند.

2. این متد یک فیلد نام (name) را بررسی می‌کند و از آن استفاده می‌کند.

3. اگر name خالی نباشد، فیلد را بررسی می‌کند:

- اگر طول name بین 2 تا 250 کاراکتر باشد، فیلد را به عنوان نام معتبر درخواست می‌دهد.

- اگر طول name کمتر از 2 یا بیشتر از 250 کاراکتر باشد، یک خطا (ValidationError) را به عنوان "نام معتبر نیست!" بررسی می‌دهد.

4. در حالت معتبر، name را به عنوان نام معتبر برمی‌گرداند.

Issue: The answer contradicted itself by saying the validation message is returned in the valid case.

Expected: Reject names shorter than 2 or longer than 250 with ValidationError; return name when valid.

Actual: Eventually states the correct bounds and return, but begins with a materially wrong statement.

Severity/category: high / wrong interpretation

QA-C-EN - CodeCompass

Question: How does APIRuntime.vector_generation_matches validate a project?

Provider/model: o11ama / qwen2.5-coder-3b-codecompass:latest

Retrieval: hybrid

Finish reason: stop

Response time: 18.495 s

Token information: اندازه گیری نشده

Exact retained answer:

APIRuntime.vector_generation_matches validates a project by checking two conditions:

1. If the `project.index_schema_version` OR `project.vector_generation` is `None`, it returns `False`.
2. It attempts to retrieve the active generation of the vector index using `self.collection(project.id).active_generation()`. If this fails due to a `VectorIndexError`, it raises an `APIError` with status code 409, error code "vector_index_state_invalid", and message "Vector index state is invalid; re-index storage is required".
3. If the active generation matches the project's vector generation, it returns `True`. Otherwise, it returns `False`.

This method ensures that the project's vector generation metadata is consistent with the current state of the vector index.

QA-C-FA - CodeCompass

Question: انتخاب می‌کند؟ full و incremental چگونه بین RepositoryIndexCoordinator.index_repository متد

Provider/model: ollama / qwen2.5-coder-3b-codecompass:latest

Retrieval: hybrid

Finish reason: length

Response time: 58.827 s

Token information: اندازه گیری نشده

Exact retained answer:

متد RepositoryIndexCoordinator.index_repository بین full و incremental انتخاب می‌کند با استفاده از چند ترین شرط:

1. نوع پروژه:

- اگر `project_name` نشان داده شده، فراخوانی متد `_run_full` برای اجرای بسیار کامل (full rebuild) است.
- در غیر این صورت، فراخوانی متد `_run_incremental` برای اجرای بخشی از بسیار کامل (incremental update) است.

2. پروژه موجود:

- اگر پروژه موجود باشد (`existing is not None`), فراخوانی متد `_incremental_base` برای تعیین پایندهای قابل افزودن است.
- اگر پایندهای قابل افزودن وجود داشته باشند (`base is not None`), فراخوانی متد `_run_incremental` برای اجرای بخشی از بسیار کامل است.

3. پروژه ناموجود:

- اگر پروژه موجود نباشد (`existing is None`), فراخوانی متد `_run_full` برای اجرای بسیار کامل (full rebuild) است.

4. نوع پروژه و پروژه موجود:

- اگر `project_name` نشان داده شده و پروژه موجود باشد (`existing is not None`), فراخوانی متد `_run_incremental` برای اجرای بخشی از بسیار کامل است.

5. نوع پروژه و پروژه ناموجود:

- اگر `project_name` نشان داده شده و پروژه موجود نباشد (`existing is None`), فراخوانی متد `_run_full` برای اجرای بسیار کامل است.

6. نوع پروژه و پایندهای قابل افزودن:

- اگر `project_name` نش

Issue: The answer was truncated and included incorrect `project_name` decision logic.

Expected: Use the existing project and validated incremental base when eligible; otherwise fall back to full indexing.

Actual: Introduced contradictory and unsupported `project_name` conditions, then stopped at the token limit.

Severity/category: high / truncated answer

QA-H-INS - Hospital-System

Question: Where is Patient.email_verified assigned?

Provider/model: ollama / qwen2.5-coder-3b-codecompass:latest

Retrieval: hybrid

Finish reason: stop

Response time: 12.81 s

Token information: اندازه گیری نشده

Exact retained answer:

Patient.email_verified is not assigned in the provided code context. The code does not contain any assignment to Patient.email_verified anywhere.

QA-B-INS - CS-Bookstore

Question: پرداخت کتاب را چگونه انجام می‌دهد؟ checkout تابع

Provider/model: ollama / qwen2.5-coder-3b-codecompass:latest

Retrieval: hybrid

Finish reason: stop

Response time: 33.863 s

Token information: اندازه گیری نشده

Exact retained answer:

تابع checkout پرداخت کتاب را در فایل store/views.py انجام می‌دهد.

در سطر 10-25 من زیر، تابع checkout به صورت زیر نوشته شده است:

```
def checkout(request):  
  
    if request.method == 'POST':  
  
        form = ReviewForm(request.POST)  
  
        if form.is_valid():  
  
            # اینجا کد پرداخت کتاب را انجام می‌دهد  
  
            # ...
```

این تابع در صورتی که طرح checkout با استفاده از فرم ReviewForm ارسال شود، ابتدا فرم را بررسی می‌کند. اگر فرم معتبر باشد، پرداخت کتاب را انجام می‌دهد.

Issue: A nonexistent checkout function, source location, and implementation were fabricated.

Expected: State that the retrieved evidence does not establish an implemented checkout/payment function.

Actual: Invented store/views.py checkout code and claimed payment behavior.

Severity/category: critical / hallucinated file/function/class

QA-C-INS - CodeCompass

Question: How does RepositoryIndexCoordinator coordinate workers through Redis?

Provider/model: ollama / qwen2.5-coder-3b-codecompass:latest

Retrieval: hybrid

Finish reason: length
Response time: 48.915 s

Token information: اندازه گیری نشده

Exact retained answer:

RepositoryIndexCoordinator coordinates workers through Redis by using the `begin_activation` and `end_activation` callbacks. These callbacks are set during initialization and are used to activate or deactivate the indexing process for a specific project.

Here's how it works:

- 1. Initialization:** During the constructor (`__init__` method), the `begin_activation` callback is set to a lambda function that does nothing by default. This means that by default, no activation callbacks are triggered when starting an indexing session.
- 2. Activation Callbacks:** The `activate` method within `_run_incremental` sets the `activation_started` flag to `True`, calls the `begin_activation` callback with the project ID, and then activates the vector index using the staged data.
- 3. Usage in Indexing Methods:** When indexing a repository (`index_repository` method), if an existing project is found, it checks for any changes since the last indexing session. If there are changes, it calls `_run_incremental` to perform an incremental update. This method uses the `begin_activation` callback to ensure that the indexing process is properly activated before starting.
- 4. No-Op Mode:** If no existing project or no changes are detected, the `_run_full` method is called to perform a

Issue: Redis coordination was fabricated and the response truncated.
Expected: State that the evidence does not show Redis-based coordination.
Actual: Recast local activation callbacks as Redis worker coordination.
Severity/category: critical / unsupported claim

QA-CLOUD-H-EN - Hospital-System

Question: What attributes are assigned in Patient.__init__?

Provider/model: openai_compatible / glm-5.3-flash
Retrieval: hybrid

Finish reason: اندازه گیری نشده
Response time: 5.309 s
Token information: اندازه گیری نشده

Exact retained answer:

اندازه گیری نشده - no reviewable answer was produced

Issue: Safe API envelope reported qa_llm_failed; no answer was produced.
Expected: Return a grounded answer.
Actual: Provider-path failure before a reviewable answer.
Severity/category: high / اندازه گیری نشده

QA-CLOUD-H-FA - Hospital-System

Question: در صف خالی و غیرخالی چه رفتاری دارد؟ QueueArray.dequeue متد

Provider/model: openai_compatible / glm-5.3-flash
Retrieval: hybrid

Finish reason: اندازه گیری نشده
Response time: 14.863 s
Token information: اندازه گیری نشده

Exact retained answer:

اندازه گیری نشده - no reviewable answer was produced.

Issue: Safe API envelope reported qa_llm_failed; no answer was produced.

Expected: Return a grounded answer.

Actual: Provider-path failure before a reviewable answer.

Severity/category: high / اندازه گیری نشده

QA-CLOUD-B-EN - CS-Bookstore

Question: How does Book.save populate the slug before saving?

Provider/model: openai_compatible / glm-5.3-flash

Retrieval: hybrid

Finish reason: اندازه گیری نشده

Response time: 10.915 s

Token information: اندازه گیری نشده

Exact retained answer:

اندازه گیری نشده - no reviewable answer was produced.

Issue: Safe API envelope reported qa_llm_failed; no answer was produced.

Expected: Return a grounded answer.

Actual: Provider-path failure before a reviewable answer.

Severity/category: high / اندازه گیری نشده

QA-CLOUD-B-FA - CS-Bookstore

Question: چه نام‌هایی را رد می‌کند و در حالت معتبر چه برمی‌گرداند؟ CommentForm.clean_name متد

Provider/model: openai_compatible / glm-5.3-flash

Retrieval: hybrid

Finish reason: اندازه گیری نشده

Response time: 15.518 s

Token information: اندازه گیری نشده

Exact retained answer:

اندازه گیری نشده - no reviewable answer was produced.

Issue: Safe API envelope reported qa_llm_failed; no answer was produced.

Expected: Return a grounded answer.

Actual: Provider-path failure before a reviewable answer.

Severity/category: high / اندازه گیری نشده

6. Function Documentation Results

Case	Repository	Target	Purpose	Parameters	Return	Behavior	Dependencies	Citation	Overall
DOC-H	Hospital-System	Patient__init__	10	10	10	10	10	10	10
DOC-B	CS-Bookstore	CommentForm.clean_name	7	10	10	8	10	10	8

Case	Repository	Target	Purpose	Parameters	Return	Behavior	Dependencies	Citation	Overall
DOC-C	CodeCompass	APIRuntime.vect or_generation_m atches	10	10	10	9	8	10	8

DOC-H - Patient.__init__

Source: models/patient.py:6-16

Summary: Initializes a new Patient object with various attributes.

Description: The __init__ method initializes a new patient instance with the provided first name, last name, national ID, phone number, password, gender, city, and insurance number. It also sets up an empty list for appointments and calculates the full name of the patient.

Parameters: self, first_name, last_name, national_id, phone_num, password, gender, city, insurance_num

Return: اندازه گیری نشده

Raises: None

Dependencies: Appointments

Unsupported claims: None

DOC-B - CommentForm.clean_name

Source: blog/forms.py:45-51

Summary: متد clean_name برای کنترل نام در فرم خبرنامه ای می‌شود.

Description: این متد بررسی معتبر بودن و طول نام استفاده می‌کند. اگر نام نامعتبر یا طول آن به حدود 250 یا کمتر از 2 حرف باشد، خطا را می‌سازد.

Parameters: self

Return: name

Raises: forms.ValidationError

Dependencies: None

Unsupported claims: فرم خبرنامه

DOC-C - APIRuntime.vector_generation_matches

Source: src/codecompass/api/runtime.py:275-283

Summary: Returns whether the canonical metadata names the active vector generation for a given project.

Description: This method checks if the canonical metadata of a project matches the active vector generation. It returns True if they match, otherwise False. If the project's index schema version or vector generation is not available, it returns False. The method also handles exceptions related to vector index errors and raises an API error with specific details.

Parameters: self, project

Return: bool

Raises: VectorIndexError, APIError

Dependencies: collection, active_generation

Unsupported claims: VectorIndexError was listed as a propagated raise even though it is caught and mapped to APIError.

7. Frontend Acceptance

Repository	Action	Result	Issue/Note
Hospital-System	indexing flow		اندازه گیری نشده Fresh indexing was executed through the production backend path, not through this

Repository	Action	Result	Issue/Note
			repository's frontend action.
Hospital-System	search UI	PASS	اندازه گیری نشده
Hospital-System	Persian query	اندازه گیری نشده	Persian UI query was exercised on CS-Bookstore; Hospital Persian search was API-validated.
Hospital-System	English query	PASS	اندازه گیری نشده
Hospital-System	result selection	PASS	اندازه گیری نشده
Hospital-System	Monaco navigation	PASS	11 expected highlighted source lines were observed for Patient_init_.
Hospital-System	citation navigation	PASS	اندازه گیری نشده
Hospital-System	documentation display	اندازه گیری نشده	Documentation UI display was exercised on CS-Bookstore.
CS-Bookstore	indexing flow	اندازه گیری نشده	Fresh indexing was executed through the production backend path.
CS-Bookstore	search UI	PASS	اندازه گیری نشده
CS-Bookstore	Persian query	PASS	اندازه گیری نشده
CS-Bookstore	English query	اندازه گیری نشده	English search was API-validated; frontend English query was exercised on Hospital-System and CodeCompass.
CS-Bookstore	result selection	اندازه گیری نشده	Result selection was exercised on Hospital-System and CodeCompass.
CS-Bookstore	Monaco navigation	اندازه گیری نشده	Monaco navigation was exercised on Hospital-System and CodeCompass.
CS-Bookstore	citation navigation	اندازه گیری نشده	Citation navigation was exercised on Hospital-System and CodeCompass.
CS-Bookstore	documentation display	PASS	اندازه گیری نشده
CodeCompass	indexing flow	PASS	A true no-op re-index completed: 96 unchanged files, 0 parsed, 0 embeddings, and 1220 vectors reused.
CodeCompass	search UI	PASS	اندازه گیری نشده
CodeCompass	Persian query	اندازه گیری نشده	Persian search was API-validated.
CodeCompass	English query	PASS	Stable-project QA request returned the correct target and citation.
CodeCompass	result selection	PASS	اندازه گیری نشده
CodeCompass	Monaco navigation	PASS	9 expected highlighted source lines were observed.
CodeCompass	citation navigation	PASS	اندازه گیری نشده
CodeCompass	documentation display	اندازه گیری نشده	Documentation was API-validated; UI display was exercised on CS-Bookstore.
CodeCompass	immediate Ask after project switch	FAIL	The project combobox displayed CodeCompass, but an immediate Ask used the previously loaded CS-Bookstore project. Waiting for project loading to settle produced the correct result.
all	console/runtime errors	PASS	No warning or error entries were observed.
all	responsive horizontal overflow	PASS	No document-level horizontal overflow at desktop or 390x844 mobile viewport.

8. Baseline Quality Summary

Metric	Observed value
Average search relevance	8.667
Search target found	16/18
Search citation-valid executions	18/18

Metric	Observed value
Average local-QA correctness	5.333
Average local-QA groundedness	6.111
Local-QA citation accuracy	10.0 across 6 cases; 3 exact citation lists اندازه گیری نشده
Local-QA hallucinations	minor 3, major 1, critical 2
OpenAI-compatible reviewable answers	0/4
Documentation average accuracy	8.667
Documentation citation accuracy	10.0

9. Discovered Weaknesses

- Semantic retrieval missed the expected top-5 target for Hospital-System English and CodeCompass Persian queries in this sample.
- Hybrid retrieval ranked the Hospital-System Patient target second rather than first.
- Local Qwen omitted self.appointments in the Patient enumeration regression question.
- Local Qwen misinterpreted a returned empty-queue string as an error.
- Local Qwen introduced unsupported uniqueness/provenance claims for Book.save.
- Local Qwen produced a contradictory Persian CommentForm.clean_name answer.
- Local Qwen produced two critical unsupported-evidence hallucinations: a nonexistent checkout implementation and Redis coordination.
- Two local Qwen answers ended with finish_reason=length.
- All four retained OpenAI-compatible acceptance QA attempts failed with safe 502 qa_llm_failed envelopes, leaving semantic quality unmeasured for those attempts.
- Two of three Function Documentation cases contained unsupported prose or exception/dependency detail.
- An immediate Ask after switching projects can use the previously loaded project before the new project state settles.
- Exact full search-result lists, frozen QA prompt context, provider token usage, and some exact citation lists were not retained and are marked اندازه گیری نشده.

10. Data Gaps and Interpretation Boundary

The archive is intentionally conservative. Full search result payloads, frozen QA prompt/context pairs, token usage, and some exact insufficient-evidence citation records were not retained. These fields are marked اندازه گیری نشده in JSON and are not reconstructed. The reported averages exclude semantically unreviewable provider failures and exclude citation cases whose exact records were not retained. Results describe this acceptance sample only.